

Text Analytics Case study

Fake news detection using Neural Net



Tejas Shinde (37) & Aditya Pandit (21)

Submitted on : 16th September, 2019

Introduction

Fake news detection is a hard problem in deception detection, and it has super real-global political and social impacts. However, statistical approaches to preventing fake news has been dramatically constrained by way of the shortage of classified benchmark information units. In this paper, we gift a version: a new, publicly to be had information set for fake news detection. We accumulated a decade-lengthy, 12.8K manually categorized brief statements in diverse contexts from POLITIFACT.COM, which offers certain analysis document and hyperlinks to source files for each case. This data set may be used for truth-checking re-seek as properly. Notably, this new facts set is an order of significance larger than previously largest public fake news information sets of similar type. Empirically, we investigate automatic faux information detection primarily based on floor-stage linguistic styles. We have designed a novel, hybrid CNN to integrate meta-data with textual content. We show that this hybrid

approach can enhance a text-handiest deep learning model.

In the past election of the 45th President of the United States, there have been many fake news post elections regarding the new President. These fake news may pose various threats to journalism, as journalists use them to reach out to every corner of the world. Last year, a man carried an AR-15 rifle and walked in a Washington DC Pizzeria, because herecently read online that “this pizzeria was harbor-ing young children as sex slaves as part of a child-abuse ring led by Hillary Clinton”¹. The man was later arrested by cops, and he was charged for firing an assault rifle in the restaurant (Kang and Goldman, 2016).

The problem of detecting fake news is quite tough as the language used in political news and financial articles are unpredictable, without applying deep learning and statistical analysis methods. The TV interviews are converted into text and treated as data for the model, but this data may contain slang as it a spoken one.

To address the issue of less data, we appeal to introduce this model, which

handles 12836 short statements labeled for truthfulness, subject, context, speaker, state, prior history.

What the data looks like

The dataset is divided into three splits, train, test and dev splits. Here are some statistics for the same.

Dataset Statistics	
Training set size	10,269
Validation set size	1,284
Testing set size	1,283
Avg. statement length (tokens)	17.9
Top-3 Speaker Affiliations	
Democrats	4,150
Republicans	5,687
None (e.g., FB posts)	2,185

In particular, PolitiFact covers a wide-range of political topics, and they provide detailed judgments with fine-grained labels. Recently, Ferreira and Vlachos (2016) have released the Emergent dataset, which includes 300 labeled rumors from PolitiFact. However, with less than a thousand samples, it is impractical to use these datasets as benchmarks for developing and evaluating machine learning algorithms for fake news detection. A sample looks like the following:

Statement: *“Newly Elected Republican Senators Sign Pledge to Eliminate Food Stamp Program in 2015.”*

Speaker: Facebook posts

Context: social media posting

Label: Pants on Fire

Justification: More than 115,000 social media users passed along a story headlined, “Newly Elected Republican Senators Sign Pledge to Eliminate Food Stamp Program in 2015.” But they failed to do due diligence and were snookered, since the story came from a publication that bills itself (quietly) as a “satirical, parody website.” We rate the claim Pants on Fire.

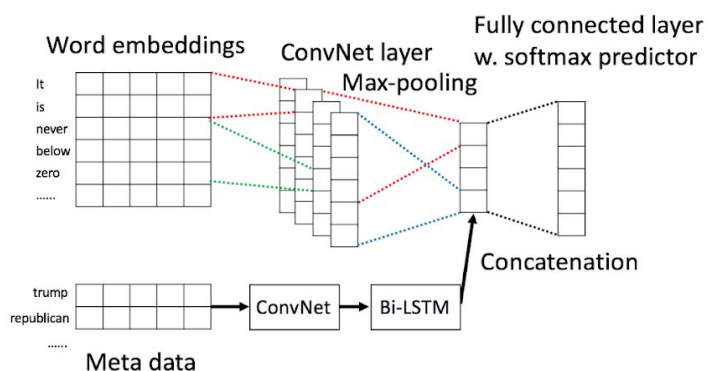
Methodology

One of the applications is to empower development of machine learning models for automatic news detection. We need to answer these questions:

1. How well can a traditional machine learning algorithm, classify a short statement into a fine grained probability value of fakeness.
2. How can we use a deep learning model, to smoothen the results as compared to the traditional machine learning techniques.

Since Convolutional Neural Networks have obtained innumerable state of the art results, we have proposed to use a

Convolutional Neural Network. We randomly initialize embedding's to encode metadata embeddings. The CNN will capture the dependence between the meta data vectors. Then, a standard max-pooling operation is performed on the latent space, followed by a bi-directional LSTM layer. We then concatenate the max-pooled text representations with the meta-data representation from the bi-directional LSTM, and feed them to fully connected layer with a softmax activation.



Whenever the model goes wrong, L2 Regularization is used to penalize the loss function, to slide towards the minima in the gradient.

In addition to features related directly to the content of the news articles, additional social context features can also be derived from the user-driven social engagements of news consumption on social media platform. Social engagements represent the news proliferation process over time, which provides useful auxiliary information

to infer the veracity of news articles. Note that few papers exist in the literature that detect fake news using social context features. However, because we believe this is a critical aspect of successful fake news detection, we introduce a set of common features utilized in similar research areas, such as rumor veracity classification on social media. Generally, there are three major aspects of the social media context that we want to represent: users, generated posts, and networks. Below, we investigate how we can extract and represent social context features from these three aspects to support fake news detection.

Evaluation Metrics

To examine the performance of algorithms for fake news detection hassle, numerous assessment metrics were used. In this subsection, we overview the most extensively used metrics for faux news detection. Most existing strategies take into account the faux information hassle as a class problem.

The Receiver Operating Characteristics (ROC) curve provides a way of comparing the performance of classifiers by looking at the trade-off in the False Positive Rate (FPR) and the True Positive

Rate(TPR). To draw the ROC curve, we plot the FPR on the x axis and and TPR along the y axis. The ROC curve compares the performance of different classifiers by changing class distributions via a threshold.

Truth discovery is the problem of detecting true data from more than one conflicting sources. Truth discovery methods do no longer explore the truth claims directly, but rely on a set of contradicting resources that document the homes of gadgets to decide the truth price. Truth discovery targets to determine the supply credibility and item truthfulness at the same time. The faux news detection hassle can benefit from numerous factors of fact discovery strategies under different situations. First, the credibility of various news stores may be modeled to infer the truthfulness of reported information. Second, applicable social media posts also can be modeled as social reaction sources to higher decide the truthfulness of claims. However, there are some other issues that ought to be considered to apply reality discovery to fake information detection in social media scenarios. First, most current truth discovery techniques consciousness on coping with structured enter within the form of Subject-Predicate-Object tuples, even as social media information is pretty unstructured and noisy.

Whenever the model goes wrong, L2 Regularization is used to penalize the loss function, to slide towards the minima in the gradient.

In addition to features related directly to the content of the news articles, additional social context features can also be derived from the user-driven social engagements of news consumption on social media platform. Social engagements represent the news proliferation process over time, which provides useful auxiliary information to infer the veracity of news articles. Note that few papers exist in the literature that detect fake news using social context features. However, because we believe this is a critical aspect of successful fake news detection, we introduce a set of common features utilized in similar research areas, such as rumor veracity classification on social media. Generally, there are three major aspects of the social media context that we want to represent: users, generated posts, and networks. Below, we investigate how we can extract and represent social context features from these three aspects to support fake news detection.

Conclusion

Compared to prior statistics units, LIAR is an order of a value larger, which enables the development

of statistical and computational techniques to faux news detection, actual-world quick statements from various contexts with various audio system additionally make there seek on growing extensive-insurance fake news detector viable. We show that while combining meta-facts with textual content, vast development test be carried out for exceptional-grained faux news detection. Given the special evaluation report and hyperlinks to supply files on this statistics set, it's also viable to explore the undertaking of automatic truth-checking over information base inside the destiny. Our corpus also can be used for stance classification, argument mining, topic modeling, rumor detection, and political NLP studies.

References

Conroy, N. J., Rubin, V. L., & Chen, Y. (2015). Automatic deception detection: Methods for finding fake news. *Proceedings of the Association for Information Science and Technology*, 52(1), 1-4.

Wang, W. Y. (2017). "liar, liar pants on fire": A new benchmark dataset for fake news detection. *arXiv preprint arXiv:1705.00648*.

